



Transformation Processing Smackdown

Spark vs Hive vs Pig

Lester Martin

DevNexus 2017

Connection before Content

Lester Martin – Hadoop/Spark Trainer & Consultant

lester.martin@gmail.com

<http://lester.website> (*links to blog, twitter, github, LI, FB, etc*)



Agenda

- **Present Frameworks**
- **File Formats**
- **Source to Target Mappings**
- **Data Quality**
- **Data Profiling**
- **Core Processing Functionality**
- **Custom Business Logic**
- **Mutable Data Concerns**
- **Performance**

! DS/ML

**Lots of
Code !**

Standard Disclaimers Apply

Wide Topic – Multiple Frameworks – Limited Time, so...

- **Simple use cases**

- Glad to enhance <https://github.com/lestermartin/oss-transform-processing-comparison>

- **ALWAYS 2+ ways to skin a cat**

- Especially with Spark ;-)

- **CLI, not GUI, tools**

- Others in that space such as Talend, Informatica & Syncsort

- **ALL code compiles in PPT ;-)**

- ***Won't explain all examples!!!***



Apache Pig – <http://pig.apache.org>

- A high-level data-flow scripting language (Pig Latin)
- Run as standalone scripts or use the interactive shell
- Executes on Hadoop
- Uses lazy execution



Grunt shell

```
rich — root@sandbox:~ — ssh — 59x5
grunt> employees = LOAD 'pigdemo.txt' AS (state, name);
grunt> describe employees;
employees: {state: bytearray,name: bytearray}
grunt> employees_grp = group employees by state;
grunt> dump employees;
```



Simple and Novel Commands

Pig Command	Description
LOAD	Read data from file system
STORE	Write data to file system
FOREACH	Apply expression to each record and output 1+ records
FILTER	Apply predicate and remove records that do not return true
GROUP/COGROUP	Collect records with the same key from one or more inputs
JOIN	Joint 2+ inputs based on a key; various join algorithms exist
ORDER	Sort records based on a key
DISTINCT	Remove duplicate records
UNION	Merge two data sets
SPLIT	Split data into 2+ more sets based on filter conditions
STREAM	Send all records through a user provided executable
SAMPLE	Read a random sample of the data
LIMIT	Limit the number of records

Executing Scripts in Ambari Pig View

The screenshot displays the Ambari Pig View interface. At the top, the navigation bar includes the Ambari logo, 'Sandbox' status with '0 ops' and '0 alerts', and links to 'Dashboard', 'Services', 'Hosts', and 'Alerts'. A user profile 'maria_dev' is logged in. On the left, a sidebar shows a file icon, a 'baseball' script, and a menu with 'Save', 'Copy', and 'Delete' options. The main area is titled 'Script' and 'History'. The 'baseball' script is selected, showing its path as '/tmp/.pigscripits/baseball-2016-03-14_10-50.pig'. The script content is as follows:

```
1 batting = load 'baseball/Batting.csv' using PigStorage(',')
2 AS (playerID:chararray, year:int, dollar2:chararray, dollar3:chararray, dollar4:chararray,
3     dollar5:chararray, dollar6:chararray, dollar7:chararray, runs:int);
4
5 raw_runs = FILTER batting BY (year > 0) AND (runs > 0);
6
7 runs = FOREACH raw_runs GENERATE playerID, year, runs;
8 grp_data = GROUP runs by (year);
9
10 max_runs = FOREACH grp_data {
11     inner_sorted = ORDER runs BY runs DESC;
12     first_row = LIMIT inner_sorted 1;
13     --GENERATE group AS grp, first_row AS the_first_row;
14     GENERATE first_row AS most_hits;
15 }
16 dump max_runs;
```

At the top right of the script editor, there is a checkbox for 'Execute on Tez' and a blue 'Execute' button.

Apache Hive – <http://hive.apache.org>

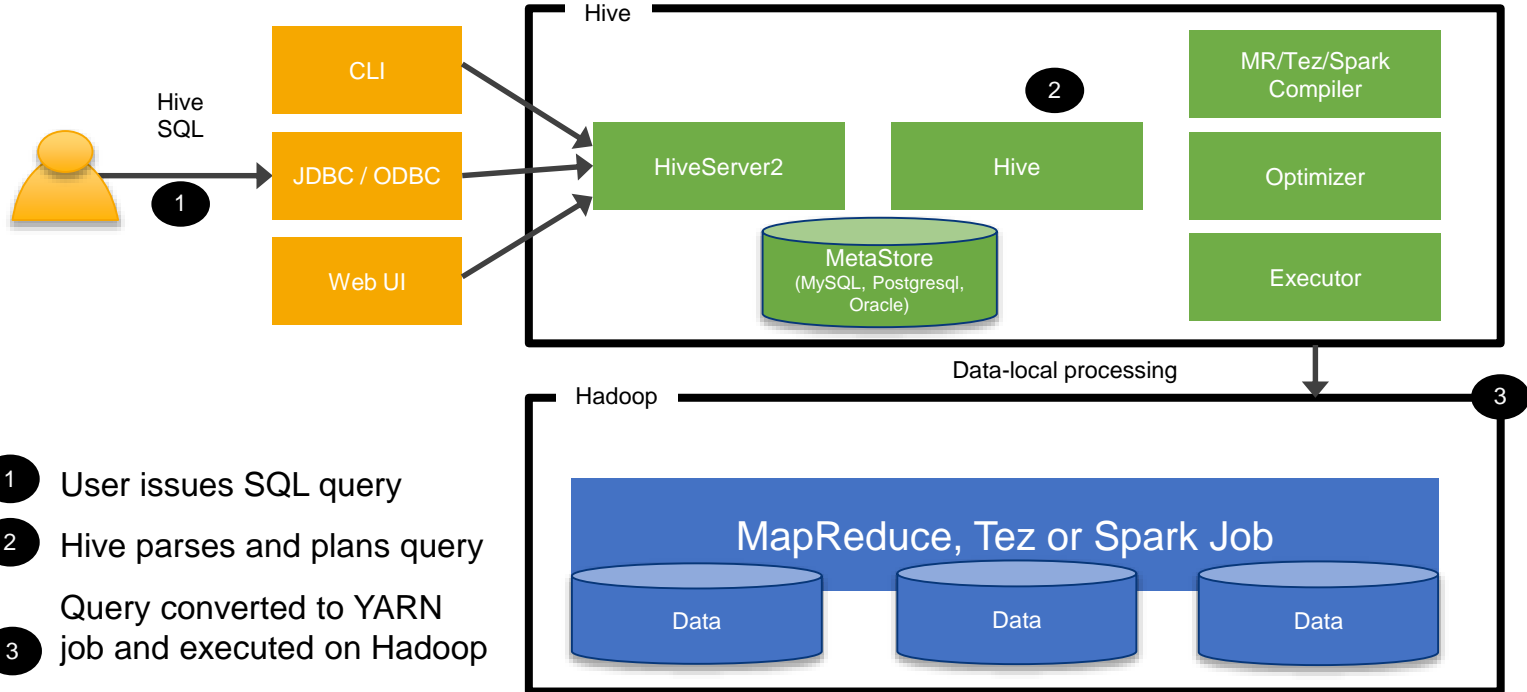
- **Data warehouse system for Hadoop**
- **Create schema/table definitions that point to data in HDFS**
- **Treat your data in Hadoop as tables**
- **SQL 92**
- **Interactive queries at scale**



Hive's Alignment with SQL

SQL Datatypes	SQL Semantics
INT	SELECT, LOAD, INSERT from query
TINYINT/SMALLINT/BIGINT	Expressions in WHERE and HAVING
BOOLEAN	GROUP BY, ORDER BY, SORT BY
FLOAT	CLUSTER BY, DISTRIBUTE BY
DOUBLE	Sub-queries in FROM clause
STRING	GROUP BY, ORDER BY
BINARY	ROLLUP and CUBE
TIMESTAMP	UNION
ARRAY, MAP, STRUCT, UNION	LEFT, RIGHT and FULL INNER/OUTER JOIN
DECIMAL	CROSS JOIN, LEFT SEMI JOIN
CHAR	Windowing functions (OVER, RANK, etc.)
VARCHAR	Sub-queries for IN/NOT IN, HAVING
DATE	EXISTS / NOT EXISTS
	INTERSECT, EXCEPT

Hive Query Process



- 1 User issues SQL query
- 2 Hive parses and plans query
Query converted to YARN job and executed on Hadoop
- 3

Submitting Hive Queries – *CLI and GUI Tools*

```
[t1@sandbox ~]$ beeline -u jdbc:hive2://localhost:10000
WARNING: Use "yarn jar" to launch YARN applications.
Connecting to jdbc:hive2://localhost:10000
Connected to: Apache Hive (version 1.2.1000.2.4.0.0-169)
Driver: Hive JDBC (version 1.2.1000.2.4.0.0-169)
Transaction isolation: TRANSACTION_REPEATABLE_READ
Beeline version 1.2.1000.2.4.0.0-169 by Apache Hive
0: jdbc:hive2://localhost:10000> show tables;
+-----+
| tab_name |
+-----+
| avg_mileage |
| driver_mileage |
| finalresults |
| geo_normal_event |
| geolocation |
| geolocation_stage |
| hcatsmokeid000a0f02_date250116 |
| risk_factor |
| risk_factor_spark |
| sample_07 |
| sample_08 |
| truck_mileage |
| trucks |
| trucks_stage |
+-----+
14 rows selected (0.696 seconds)
0: jdbc:hive2://localhost:10000> 
```



Submitting Hive Queries – *Ambari Hive View*

The screenshot displays the Ambari Hive View interface. At the top, the navigation bar includes 'Ambari', 'Sandbox', '0 ops', '0 alerts', 'Dashboard', 'Services', 'Hosts', 'Alerts', and a user profile 'maria_dev'. Below this, the 'Hive' tab is selected, with sub-tabs for 'Query', 'Saved Queries', 'History', 'UDFs', and 'Upload Table'.

The main interface is divided into three sections:

- Database Explorer:** Located on the left, it shows a tree view of databases under 'default'. The 'default' database is expanded, listing tables such as 'avg_mileage', 'driver_mileage', 'finalresults', 'geo_normal_event', 'geolocation', 'geolocation_stage', 'hcatsmokeid000a0f02...', 'risk_factor', 'risk_factor_spark', 'sample_07', and 'xademo'.
- Query Editor:** The central area for writing and executing queries. It shows a worksheet named 'avgmpg' with the following SQL query:

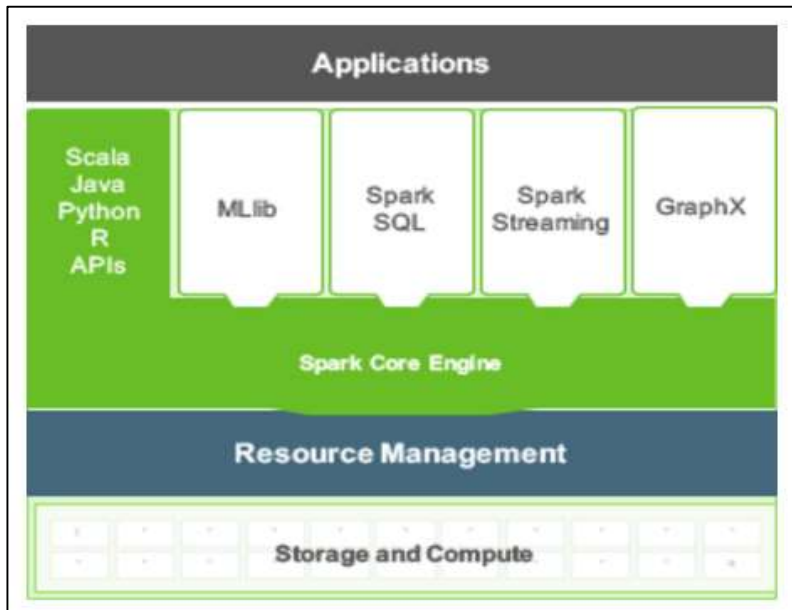
```
1 SELECT truckid, avg(mpg) avgmpg
2   FROM truck_mileage
3  GROUP BY truckid;
```

Below the query editor are buttons for 'Execute', 'Explain', 'Save as...', 'Kill Session', and 'New Worksheet'.
- Query Process Results (Status: Succeeded):** The bottom section displays the results of the executed query. It has tabs for 'Logs' and 'Results'. The 'Results' tab is active, showing a table with columns 'truckid' and 'avgmpg'. The table contains five rows of data.

On the right side of the interface, there is a vertical toolbar with icons for 'SQL', 'Settings', 'Alerts', 'TEZ', and a notification icon.

truckid	avgmpg
A1	4.785822711239916
A10	5.401717663765759
A100	4.939038953107008
A11	5.502368692859457

Apache Spark – <http://spark.apache.org>




- ◆ A data access engine for fast, large-scale data processing
- ◆ Designed for iterative in-memory computations and interactive data mining
- ◆ Provides expressive multi-language APIs for Scala, Java, R and Python
- ◆ Data workers can use built-in libraries to rapidly iterate over data for:
 - ETL
 - Machine learning
 - SQL workloads
 - Stream processing
 - Graph computations



Spark Executors & *Cluster Deployment Options*

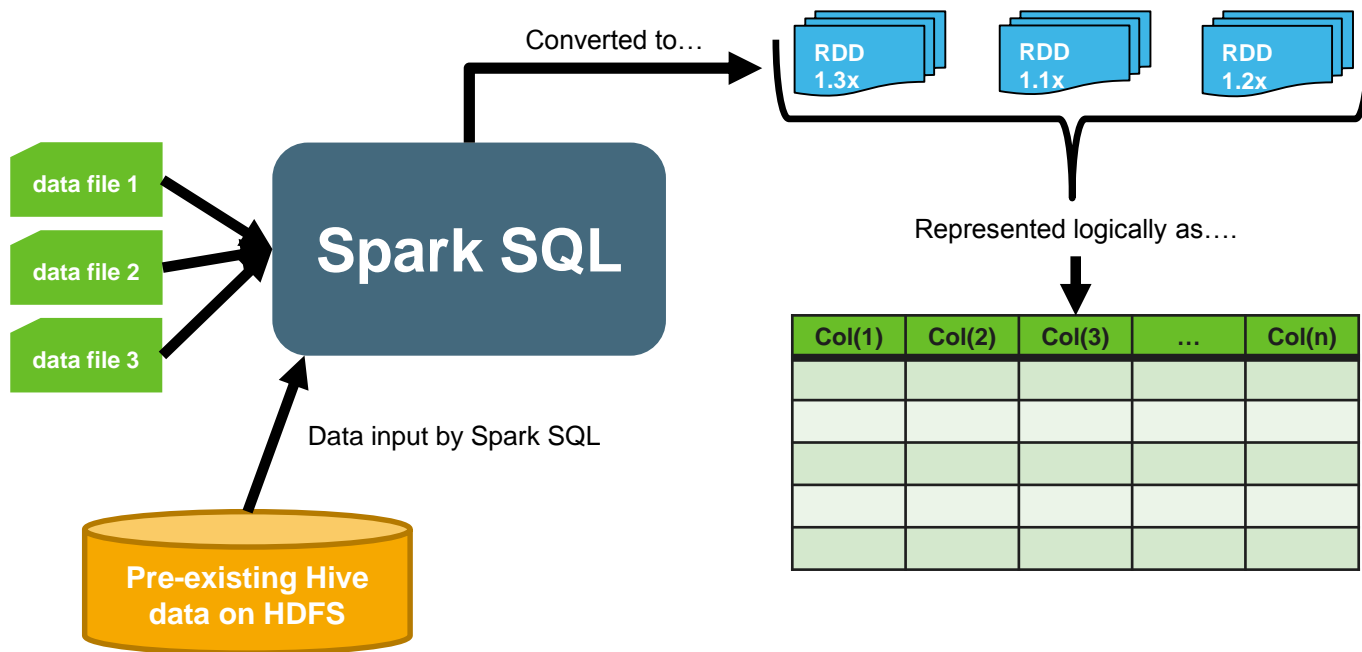
- Responsible for all application workload processing
 - The "workers" of a Spark application
 - Includes the `SparkContext` serving as the "master"
 - Schedules tasks
 - Pre-created in shells & notebooks
- Exist for the life of the application**
- Standalone mode and cluster options
 - YARN
 - Mesos



Spark SQL Overview

- ◆ A module built on top of Spark Core
- ◆ Provides a programming abstraction for distributed processing of large-scale structured data in Spark
- ◆ Data is described as a DataFrame with rows, columns and a schema
- ◆ Data manipulation and access is available with two mechanisms
 - SQL Queries
 - DataFrames API

The DataFrame Visually



Apache Zeppelin – <http://zeppelin.apache.org>

Zeppelin Notebook • Interpreter Configuration Search in your notebooks Connected

Australian Dataset (SparkSQL example)

Register RDD as table

```
case class Health (year: String, state: String, category:String, funding_src1: String, funding_src2: String, spending: Integer)
val health = dataset.map(x => x.split(",")).map(
  k => Health(k(0),k(1),k(2),k(3), k(4), k(5).toInt)
)
// toDF() works only in spark 1.5.0
// for spark 1.1.x and spark 1.2.x,
// use below instead:
// health.registerTempTable("health_table")
health.toDF().registerTempTable("health_table")

defined class Health
health: org.apache.spark.rdd.RDD[Health] = MapPartitionsRDD[7] at map at <console>:33
Task 3 complete
```

Spending (in billions) by state

```
%sql
select state, sum(spending)/1000 SpendingInBillions
from health_table
group by state
order by SpendingInBillions desc
```

Spending (In Billions) By Year

```
%sql
select year, sum(spending)/1000 SpendingInBillions
from health_table
group by year
order by SpendingInBillions
```

Spending (in billions) by area

```
%sql
select category, sum(spending)/1000 SpendingInBillions
from health_table
group by category
order by SpendingInBillions desc
```

category	SpendingInBillions
Public hospitals	445.845
Medical services	272.507
Private hospitals	121.022
Benefit-paid pharmaceuticals	104.221
Dental services	90.788
Community health	75.765
Capital expenditure	72.888
All other medications	70.506
Other health practitioners	51.382
Administration	41.029
Research	40.074
Aids and appliances	37.155
Patient transport services	36.174
Public health	27.072
Medical expense tax rebate	0.0

Still Based on MapReduce Principles

```
sc.textFile("/some-hdfs-data") \
  .flatMap(lambda line: line.split(" ")) \
  .map(lambda line: (word, 1)) \
  .reduceByKey(lambda a,b : a+b, \
    numPartition=3) \
  .collect()
```


RDD[String]

RDD[List[String]]

RDD[(String, Int)]

RDD[(String, Int)]

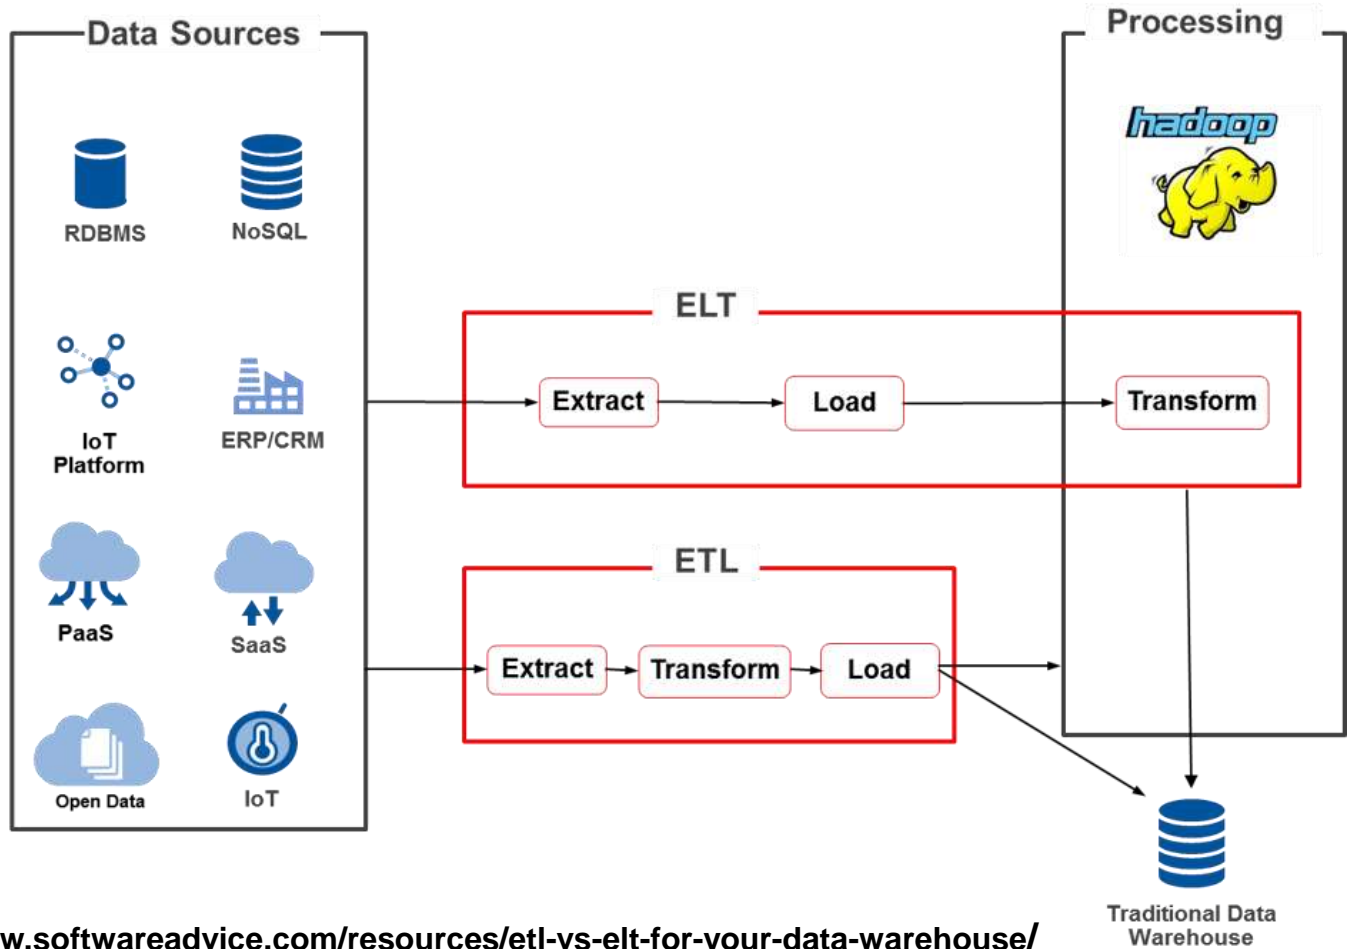
Array[(String, Int)]



ETL Requirements

- **Read/write multiple file formats and persistent stores**
- **Source-to-target mappings**
- **Data profiling**
- **Data quality**
- **Common processing functionality**
- **Custom business rules injection**
- **Merging changed records**
- **Error handling**
- **Alerts / notifications**
- **Logging**
- **Lineage & job statistics**
- **Administration**
- **Reusability**
- **Performance & scalability**
- **Source code management**

ETL vs ELT



File Formats

The ability to read & write many different file formats is critical

- **Delimited values (comma, tab, etc)**
- **XML**
- **JSON**
- **Avro**
- **Parquet**
- **ORC**
- **Esoteric formats such as EBCDIC and compact RYO solutions**

File Formats: Delimited Values

Delimited datasets are very common place in Big Data clusters

Simple example file: `catalog.del`

```
Programming Pig|Alan Gates|23.17|2016
```

```
Apache Hive Essentials|Dayong Du|39.99|2015
```

```
Spark in Action|Petar Zecevic|41.24|2016
```

Pig Code for Delimited File

```
book_catalog = LOAD
    '/otpc/ff/del/data/catalog.del'
    USING PigStorage('|')
    AS (title:chararray, author:chararray,
        price:float, year:int);

DESCRIBE book_catalog;
DUMP book_catalog;
```

Pig Output for Delimited File

```
book_catalog: {title: chararray,author:  
chararray,price: float,year: int}
```

```
(Programming Pig,Alan Gates,23.17,2016)
```

```
(Apache Hive Essentials,Dayong  
Du,39.99,2015)
```

```
(Spark in Action,Petar Zecevic,41.24,2016)
```


Hive Code for Delimited File

```
CREATE EXTERNAL TABLE book_catalog_pipe(  
    title string, author string,  
    price float, year int)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY '|'   
STORED AS TEXTFILE  
LOCATION '/otpc/ff/del/data';
```

Hive Output for Delimited File Schema

```
desc book_catalog_pipe;
```

col_name	data_type	comment	
+-----+	+-----+	+-----+	+--+
title	string		
author	string		
price	float		
year	int		

Hive Output for Delimited File Contents

```
SELECT * FROM book_catalog_pipe;
```

title	author	price	year	
+-----+	+-----+	+-----+	+-----+	+--+
Progra...	Alan...	23.17	2016	
Apache...	Dayo...	39.99	2015	
Spark ...	Peta...	41.24	2016	

Spark Code for Delimited File

```
val catalogRDD = sc.textFile(
    "hdfs:///otpc/ff/del/data/catalog.del")

case class Book(title: String, author:
String, price: Float, year: Int)

val catalogDF = catalogRDD
    .map(b => b.split('|'))
    .map(b => Book(b(0), b(1), b(2).toFloat,
        b(3).toInt))

    .toDF()
```

Spark Output for Delimited File Schema

```
catalogDF.printSchema()
```

```
root
```

```
|-- title: string (nullable = true)
|-- author: string (nullable = true)
|-- price: float (nullable = false)
|-- year: integer (nullable = false)
```

Spark Output for Delimited File Contents

```
catalogDF.show()
```

```
|          title|          author|price|year|
+-----+-----+-----+-----+
|Programming...|  Alan Gates|23.17|2016|
|Apache Hive...|   Dayong Du|39.99|2015|
|Spark in Ac...|Petar Ze...|41.24|2016|
```

File Formats: XML

Simple example file: `catalog.xml`

```
<CATALOG>
  <BOOK>
    <TITLE>Programming Pig</TITLE>
    <AUTHOR>Alan Gates</AUTHOR>
    <PRICE>23.17</PRICE>
    <YEAR>2016</YEAR>
  </BOOK>
  <!-- other 2 BOOKs not shown -->
</CATALOG>
```

Pig Code for XML File

```
raw = LOAD '/otpc/ff/xml/catalog.xml'
      USINGXMLLoader('BOOK') AS (x:chararray);

formatted = FOREACH raw GENERATE
              XPath(x, 'BOOK/TITLE') AS title,
              XPath(x, 'BOOK/AUTHOR') AS author,
              (float) XPath(x, 'BOOK/PRICE') AS price,
              (int) XPath(x, 'BOOK/YEAR') AS year;
```


Hive Code for XML File

```
CREATE EXTERNAL TABLE book_catalog_xml(str string)
LOCATION '/otpc/ff/xml/flat';
```

```
CREATE TABLE book_catalog STORED AS ORC AS

SELECT xpath_string(str, 'BOOK/TITLE') AS title,
        xpath_string(str, 'BOOK/AUTHOR') AS author,
        xpath_float( str, 'BOOK/PRICE') AS price,
        xpath_int(    str, 'BOOK/YEAR') AS year
FROM book_catalog_xml;
```

Spark Code for XML File

```
val df = sqlContext
    .read
    .format("com.databricks.spark.xml")
    .option("rowTag", "BOOK")
    .load("/otpc/ff/xml/catalog.xml")
```

File Formats: JSON

Simple example file: `catalog.json`

```
{"title": "Programming Pig", "author": "Alan Gates",  
  "price": 23.17, "year": 2016}  
{"title": "Apache Hive Essentials", "author": "Dayong Du",  
  "price": 39.99, "year": 2015}  
{"title": "Spring in Action", "author": "Petar Zecevic",  
  "price": 41.24, "year": 2016}
```

Pig Code for JSON File

```
book_catalog =  
    LOAD  '/otpc/ff/json/data/catalog.json'  
        USING  JsonLoader('title:chararray,  
                           author:chararray,  
                           price:float,  
                           year:int');
```

Hive Code for JSON File

```
CREATE EXTERNAL TABLE book_catalog_json(  
    title string, author string,  
    price float, year int)  
ROW FORMAT SERDE 'o.a.h.h.d.JsonSerDe'  
STORED AS TEXTFILE  
LOCATION '/otpc/ff/json/data';
```

Spark Code for JSON File

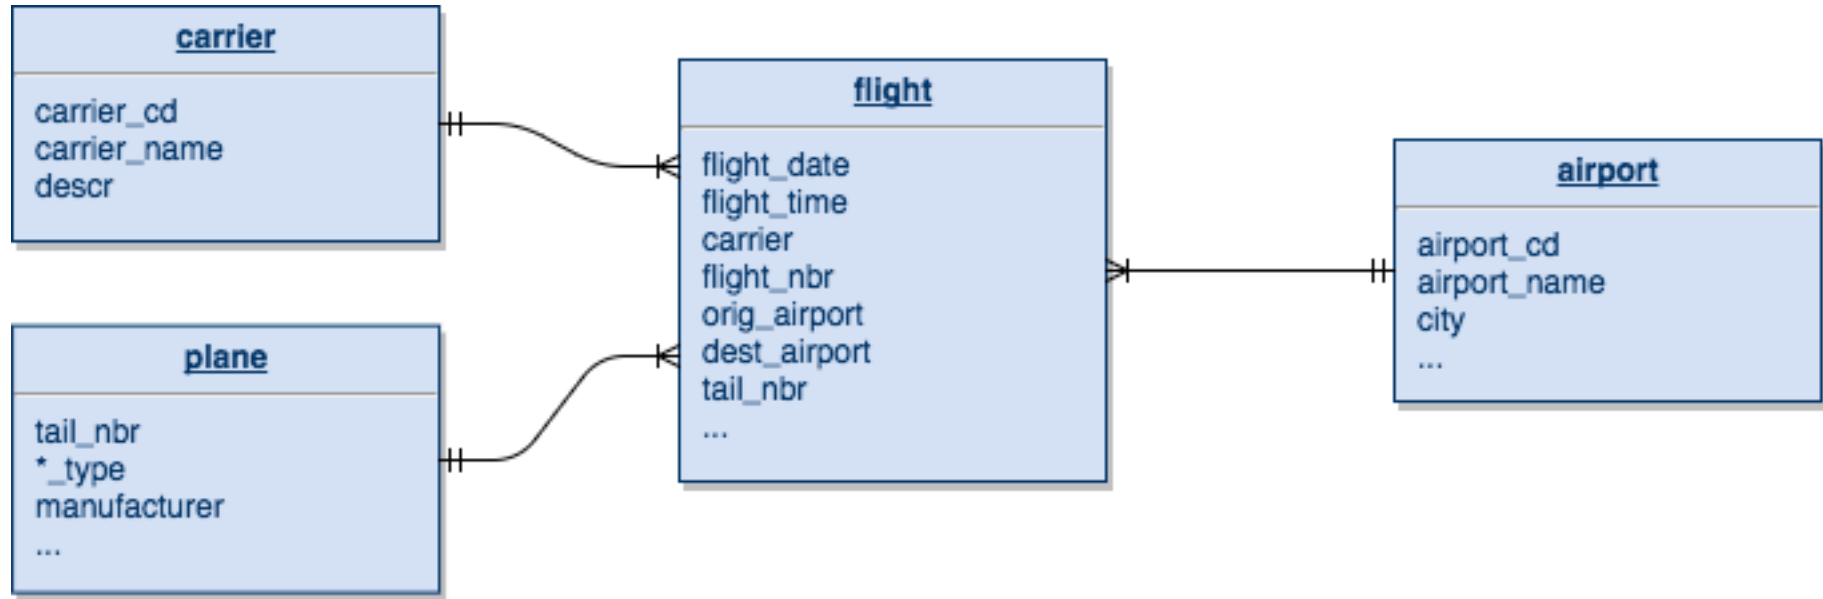
```
val df = sqlContext  
    .read  
    .format("json")  
    .load("/otpc/ff/json/catalog.json")
```

WINNER: File Formats



Data Set for Examples

Reliance on Hive Metastore for Pig and Spark SQL



Source to Target Mappings

Classic ETL need to map one dataset to another; includes these scenarios

Column Presence	Action
Source and Target	Move data from source column to target column (could be renamed, cleaned, transformed, etc)
Source, not in Target	Ignore this column
Target, not in Source	Implies a hard-coded or calculated value will be inserted or updated

Source to Target Mappings Use Case

Create new dataset from `airport_raw`

- **Change column names**
 - `airport_code` to `airport_cd`
 - `airport` to `name`
- **Carry over as named**
 - `city, state, country`
- **Exclude**
 - `latitude` and `longitude`
- **Hard-code new field**
 - `gov_agency` as `'FAA'`

Pig Code for Data Mapping

```
src_airport = LOAD 'airport_raw'
               USING o.a.h.h.p.HCatLoader();

tgt_airport = FOREACH src_airport GENERATE
               airport_code AS airport_cd,
               airport AS name, city, state, country,
               'FAA' AS gov_agency:chararray;

DESCRIBE tgt_airport;

DUMP tgt_airport;
```

Pig Output for Data Mapping

```
target_airport: {airport_cd: chararray,  
name: chararray, city: chararray, state:  
chararray, country: chararray, gov_agency:  
chararray}
```

```
(00M,Thigpen,BaySprings,MS,USA,FAA)
```

```
(00R,LivingstonMunicipal,Livingston,TX,USA,FAA)
```

```
(00V,MeadowLake,ColoradoSprings,CO,USA,FAA)
```

Hive Code for Data Mapping

```
CREATE TABLE tgt_airport STORED AS ORC AS
  SELECT airport_code AS airport_cd,
         airport AS name,
         city, state, country,
         'FAA' AS gov_agency
  FROM airport_raw;
```

Hive Output for Mapped Schema

col_name	data_type	comment	
+-----	+-----	+-----	+--+
airport_cd	string		
name	string		
city	string		
state	string		
country	string		
gov_agency	string		

Hive Output for Mapped Contents

```
SELECT * FROM tgt_airport;
```

<i>_cd</i>	<i>name</i>	<i>city</i>	<i>st</i>	<i>cny</i>	<i>g_a</i>
00M	Thig...	BayS...	MS	USA	FAA
00R	Livi...	Livi...	TX	USA	FAA
00V	Mead...	Colo...	CO	USA	FAA

Spark Code for Data Mapping

```
val airport_target = hiveContext
    .table("airport_raw")
    .drop("lat").drop("long")
    .withColumnRenamed("airport_code",
                        "airport_cd")
    .withColumnRenamed("airport", "name")
    .withColumn("gov_agency", lit("FAA"))
```


Spark Output for Mapped Schema

root

```
|-- airport_cd: string (nullable = true)
|-- name: string (nullable = true)
|-- city: string (nullable = true)
|-- state: string (nullable = true)
|-- country: string (nullable = true)
|-- gov_agency: string (nullable = false)
```

Spark Output for Mapped Contents

```
airport_target.show()
```

<code> _cd </code>	<code>name </code>	<code>city </code>	<code>st </code>	<code>cny </code>	<code>g_a </code>
<code>+---+</code>	<code>-----+</code>	<code>-----+</code>	<code>--+</code>	<code>---+</code>	<code>---+</code>
<code> 00M </code>	<code>Thigpen </code>	<code>BaySprings </code>	<code>MS </code>	<code>USA </code>	<code>FAA </code>
<code> 00R </code>	<code>Livingston...</code>	<code>Livingston...</code>	<code>TX </code>	<code>USA </code>	<code>FAA </code>
<code> 00V </code>	<code>MeadowLake </code>	<code>ColoradoSp...</code>	<code>CO </code>	<code>USA </code>	<code>FAA </code>

WINNER: Source to Target Mapping



Data Quality

DQ is focused on detecting/correcting/enhancing input data

- **Data type conversions / casting**
- **Numeric ranges**
- **Currency validation**
- **String validation**
 - Leading / trailing spaces
 - Length
 - Formatting (ex: SSN and phone #s)
- **Address validation / standardization**
- **Enrichment**

Numeric Validation Use Case

Validate `latitude` and `longitude` values from `airport_raw`

- Convert them from string to float
- Verify these values are within normal ranges

Attribute	Min	Max
<code>latitude</code>	-90	+90
<code>longitude</code>	-180	+180

Pig Code for Numeric Validation

```
src_airport = LOAD 'airport_raw' USING HCatLoader();
aprt_cnvrtd = FOREACH src_airport GENERATE
    airport_cd, (float) latitude, (float) longitude;

ll_not_null = FILTER aprt_cnvrtd BY
    ( NOT ( (latitude IS NULL) OR
            (longitude IS NULL) ) );

valid_airports = FILTER ll_not_null BY
    (latitude <= 70) AND (longitude >= -170);
```

Hive Code for Numeric Validation

```
CREATE TABLE airport_stage STORED AS ORC AS
SELECT airport_code,
        CAST(latitude AS float),
        CAST(longitude AS float)
FROM airport_raw;

CREATE TABLE airport_final STORED AS ORC AS
SELECT * FROM airport_stage
WHERE latitude BETWEEN -80 AND 70
      AND longitude BETWEEN -170 AND 180;
```

Spark Code for Numeric Validation

```
val airport_validated = hiveContext
    .table("airport_raw")
    .selectExpr("airport_code",
        "cast(latitude as float) latitude",
        "cast(longitude as float) longitude")
    .filter("latitude is not null")
    .filter("longitude is not null")
    .filter("latitude <= 70")
    .filter("longitude >= -170")
```


String Validation Use Case

Validate `city` values from `airport_raw`

- **Trim any leading / trailing spaces**
- **Truncate any characters beyond the first 30**

Pig Code for String Validation

```
src_airport = LOAD 'airport_raw'
              USING HCatLoader();

valid_airports = FOREACH src_airport
                  GENERATE
                    airport_code, airport,
                    SUBSTRING(TRIM(city),0,29) AS city,
                    state, country;
```

Hive Code for String Validation

```
CREATE TABLE airport_final STORED AS ORC AS
SELECT airport_code, airport,
       SUBSTR(TRIM(city),1,30) AS city,
       state, country
FROM airport_raw;
```

Spark Code for String Validation

```
val airport_validated = hiveContext
    .table("airport_raw")
    .withColumnRenamed("city", "city_orig")
    .withColumn("city", substring(
        trim($"city_orig"), 1, 30))
    .drop("city_orig")
```

WINNER: Data Quality



Data Profiling

Technique used to examine data for different purposes such as determining accuracy and completeness – drives DQ improvements

- **Numbers of records – including null counts**
- **Avg / max lengths**
- **Distinct values**
- **Min / max values**
- **Mean**
- **Variance**
- **Standard deviation**

Data Profiling with Pig

Coupled with Apache DataFu generates statistics such as the following

Column Name: sales_price

Row Count: 163794

Null Count: 0

Total Value: 21781793

Distinct Count: 1446

Mean Value: 132.98285040966093

Highest Value: 70589

Variance: 183789.18332067598

Lowest Value: 1

Standard Deviation: 428.7064069041609

Data Profiling with Hive

Column-level statistics for all data types

col_name	min	max	nulls	dist_ct
----------	-----	-----	-------	---------

+-----+	---+	---+	-----+	-----+
---------	------	------	--------	--------

air_time	0	757	0	316
----------	---	-----	---	-----

col_name	dist_ct	avgColLn	mxColLn
----------	---------	----------	---------

+-----+	-----+	-----+	-----+
---------	--------	--------	--------

city	2535	8.407	32
------	------	-------	----

Data Profiling with Spark

Inherent statistics for numeric data types

```
| summary |          air_time |
+-----+-----+
|  count |          2056494 |
|   mean | 103.9721783773743 |
| stddev | 67.42112792270458 |
|   min |              0 |
|   max |             757 |
```

WINNER: Data Profiling



Core Processing Functionality

Expected features to enable data transformation, cleansing & enrichment

- **Filtering / splitting**
- **Sorting**
- **Lookups / joining**
- **Union / distinct**
- **Aggregations / pivoting**
- **SQL support**
- **Analytical functions**

Filtering Examples; Pig, Hive & Spark

```
tx_arprt = FILTER arprt BY state == 'TX';
```



```
SELECT * FROM arprt WHERE state = 'TX';
```



```
val txArprt = hiveContext  
    .table("arprt")  
    .filter("state = 'TX'")
```



Sorting Examples; Pig, Hive & Spark

```
srt_flight = ORDER flight BY dep_delay DESC,  
                unique_carrier, flight_num;
```



```
SELECT * FROM flight ORDER BY dep_delay DESC,  
                unique_carrier, flight_num;
```



```
val longestDepartureDelays = hiveContext  
    .table("flight").sort($"dep_delay".desc,  
                          $"unique_carrier", $"flight_num")
```



Joining with Pig

```
jnRslt = JOIN flights BY tail_num,  
           planes BY tail_number;  
prettier = FOREACH with_year GENERATE  
    flights::flight_date AS flight_date,  
    flights::tail_num AS tail_num,  
    -- plus other 17 "flights" attribs  
    planes::year AS plane_built;  
    -- ignore other 8 "planes" attribs
```

Joining with Hive

```
SELECT F.*,  
       P.year AS plane_built  
FROM flight F,  
       plane P  
WHERE F.tail_num = P.tail_number;
```

Joining with Spark

```
val flights = hiveContext.table("flight")
val planes = hiveContext.table("plane")
    .select("tail_number", "year")
    .withColumnRenamed("year", "plane_built")
val augmented_flights = flights
    .join(planes)
    .where($"tail_num" === $"tail_number")
    .drop("tail_number")
```


Pig Code for Distinct

```
planes  = LOAD 'plane' USING HCatLoader();  
rotos = FILTER planes BY  
    aircraft_type == 'Rotorcraft';  
makers = FOREACH rotos GENERATE  
    manufacturer;  
  
distinct_makers = DISTINCT makers;  
DUMP distinct_rotor_makers;
```

Pig Output for Distinct

(BELL)

(SIKORSKY)

(AGUSTA SPA)

(AEROSPATIALE)

(COBB INTL/DBA ROTORWAY INTL IN)

Hive Code for Distinct

```
SELECT DISTINCT (manufacturer)
  FROM plane
 WHERE aircraft_type = 'Rotorcraft';
```

Hive Output for Distinct

manufacturer	
+-----+	
AEROSPATIALE	
AGUSTA SPA	
BELL	
COBB INTL/DBA ROTORWAY INTL IN	
SIKORSKY	

Spark Code for Distinct

```
val rotor_makers = hiveContext
    .table("plane")
    .filter("aircraft_type = 'Rotorcraft'")
    .select("manufacturer")
    .distinct()
```

Spark Output for Distinct

```
|          manufacturer |  
+-----+  
|          BELL |  
|      SIKORSKY |  
|    AGUSTA SPA |  
| AEROSPATIALE |  
| COBB INTL/DBA ROT... |
```

Pig Code for Aggregation

```
flights = LOAD 'flight' USING HCatLoader();
reqd_cols = FOREACH flights GENERATE
    origin, dep_delay;
by_orig = GROUP reqd_cols BY origin;
avg_delay = FOREACH by_orig GENERATE
    group AS origin,
    AVG(reqd_cols.dep_delay) AS avg_dep_delay;
srted_delay = ORDER avg_delay BY avg_dep_delay DESC;
top5_delay = LIMIT srted_delay 5;
DUMP top5_delay;
```

Pig Output for Aggregation

(PIR, 49.5)

(ACY, 35.916666666666664)

(ACK, 25.558333333333334)

(CEC, 23.40764331210191)

(LMT, 23.40268456375839)

Hive Code for Aggregation

```
SELECT origin,  
        AVG(dep_delay) AS avg_dep_delay  
FROM flight  
GROUP BY origin  
ORDER BY avg_dep_delay DESC  
LIMIT 5;
```

Hive Output for Aggregation

origin	avg_dep_delay
PIR	49.5
ACY	35.916666666666664
ACK	25.558333333333334
CEC	23.40764331210191
LMT	23.40268456375839


Spark Code for Aggregation

```
val sorted_orig_timings = hiveContext
    .table("flight")
    .select("origin", "dep_delay")
    .groupBy("origin").avg()
    .withColumnRenamed("avg(dep_delay)",
                        "avg_dep_delay")
    .sort($"avg_dep_delay".desc)
sorted_orig_timings.show(5)
```

Spark Output for Aggregation

origin	avg_dep_delay
PIR	49.5
ACY	35.916666666666664
ACK	25.558333333333334
CEC	23.40764331210191
LMT	23.40268456375839

WINNER: Core Processing Functionality



Custom Business Logic

Implemented via User Defined Functions (UDF)


- **Pig and Hive**
 - Write Java and compile to a JAR
 - Register JAR
- **Hive can administratively pre-register UDFs at the database level**
- **Spark can wrap functions at runtime**

```
+-----+-----+ from pyspark.sql.functions import udf                                +----+-----+
|      date|product| from pyspark.sql.types import IntegerType                            |year|product|
+-----+-----+
|2015-03-12|toaster| get_year = udf(lambda x: int(x[:4]), IntegerType())                |2015|toaster|
|2015-04-12|  iron| df1.select(get_year(df1["date"]).alias("year"),                    |2015|  iron|
|2014-12-31|fridge| df1.select(get_year(df1["date"]).alias("year"),                    |2014|fridge|
|2015-02-03|  cup| df1.select(get_year(df1["date"]).alias("year"),                    |2015|  cup|
+-----+-----+
                                     df1["product"])
                                     .collect()
```

WINNER: Custom Business Logic



Mutable Data – *Merge & Replace*



Ingest – bring over the incremental data

Reconcile – perform the merge

Compact – replace the existing data with the newly merged content

Purge – cleanup & prepare to repeat

See my preso and video on this topic

- <http://www.slideshare.net/lestermartin/mutable-data-in-hives-immutable-world>
- <https://www.youtube.com/watch?v=EUz6Pu1IBHQ>

WINNER: Mutable Data



Performance

Scalability is based on size of cluster


- **Tez and Spark has MR optimizations**

- Can link multiple maps and reduces together without having to write intermediate data to HDFS
- Every reducer does not require a map phase

- **Hive and Spark SQL have query optimizers**

- **Spark has the edge**

- Caching data to memory can avoid extra reads from disk
- **Resources dedicated for entire life of the application**
- Scheduling of tasks from 15-20s to 15-20ms



WINNER: Performance



Recommendations

Review ALL THREE frameworks back at “your desk”

Decision Criteria...

- Existing investments
- Forward-looking beliefs
- Adaptability & current skills
- It's a “matter of style”
- Polyglot programming is NOT a bad thing!!

Share your findings via blogs and local user groups



Questions?

Lester Martin – Hadoop/Spark Trainer & Consultant

lester.martin@gmail.com

<http://lester.website> (*links to blog, twitter, github, LI, FB, etc*)

THANKS FOR YOUR TIME!!